



Optimization & Performance



Optimization

- **Premature optimization is the root of all evil**
 - Donald Knuth (most important person in Algorithms)
 - See: https://en.wikiquote.org/wiki/Donald_Knuth
- Don't optimize until you have problems
 - Then test to make sure you fixed the right thing

Analyze Problems

- When you do **experience problems**:
 - Check the generated SQL to see what's happening

```
<properties>
  <property name="javax.persistence.jdbc.url" value="jdbc:mysql://localhost:3306/cs544?useSSL=false"/>
  <property name="javax.persistence.jdbc.driver" value="com.mysql.jdbc.Driver"/>
  <property name="javax.persistence.jdbc.user" value="root"/>
  <property name="javax.persistence.jdbc.password" value="root"/>
  <property name="hibernate.dialect" value="org.hibernate.dialect.MySQL5InnoDBDialect" />

  <!-- Useful for analyzing problems -->
  <property name="hibernate.show_sql" value="true" />
  <property name="hibernate.format_sql" value="true" />

  <!-- 2nd Level Caching -->
  <property name="hibernate.cache.provider_class" value="org.hibernate.cache.NoCacheProvider"/>
  <property name="hibernate.generate_statistics" value="true" />

  <property name="hibernate.id.new_generator_mappings" value="false" />
  <property name="hibernate.hbm2ddl.import_files" value="test.sql" />
  <property name="javax.persistence.schema-generation.database.action" value="drop-and-create"/>
</properties>
```

Optimization

- First we will talk about LAZY and EAGER loading
 - And we will describe **common problems**
- To solve problems we'll look at:
 - Standard JPA ways to solve problems
 - Hibernate extensions to solve problems



CS544 EA

Hibernate

Lazy and Eager


What is Lazy (or Eager)?

- Lazy means it does not load it
 - Until it absolutely needs it
 - By default all **Collections** (*ToMany) **are lazy**
- Eager gets loaded right away
 - As soon as it knows about it
 - By default all **references** (*ToOne) **are eager**

Changing FetchType

- It's **possible to change**
 - References to LAZY
 - collections to EAGER
- Generally not needed

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private Long id;
    private String firstName;
    private String lastName;
    @ManyToOne(fetch=FetchType.LAZY)
    private SalesRep salesRep;
    @OneToMany(fetch=FetchType.LAZY)
    @JoinColumn
    private List<Book> books
        = new ArrayList<>();
}
```



LAZY or EAGER

- Some say that all associations should be lazy
 - EAGERly loaded objects may never be used
- Usually Hibernate loads *ToOne with Joins
 - Less expensive than separate selects
 - Still takes overhead, wasted if object not used
- On the whole **you can be safe never modifying**
 - Premature optimization is the root of all evil

@LazyCollection

- A Hibernate extension that:
 - Can be useful **for big collections**
- By default the entire collection is retrieved for:
 - `.size()`, `.isEmpty()`, `.contains()`
 - Instead of **using the DB** to count / check
 - ‘Extra lazy’ fixes that

@LazyCollection

```
import org.hibernate.annotations.LazyCollection;  
import org.hibernate.annotations.LazyCollectionOption;
```

@Entity

```
public class Customer {  
    @Id  
    @GeneratedValue  
    private Long id;  
    private String firstName;  
    private String lastName;  
    @OneToMany(cascade=CascadeType.ALL)  
    @JoinColumn  
    @LazyCollection(LazyCollectionOption.EXTRA)  
    private List<Movie> movies  
        = new ArrayList<>();  
}
```

Extra Lazy

```
Customer c = em.find(Customer.class, 1L);  
System.out.println(c.getMovies().size());
```

Hibernate:

```
select  
    customer0_.id as id1_1_0_,  
    customer0_.firstName as firstNam2_1_0_,  
    customer0_.lastName as lastName3_1_0_,  
    customer0_.salesRep_id as salesRep4_1_0_  
from  
    Customer customer0_  
where  
    customer0_.id=?
```

Hibernate:

```
select  
    count(id)  
from  
    Movie  
where  
    movies_id =?
```

Lazy Properties

- It is possible to make **individual properties lazy**
 - Needs Property access (getters)
 - `@Basic(fetch=FetchType.LAZY)` (on getters)
- **Needs ByteCode instrumentation** to work
 - Rewrites your getters (after compilation) for ability to load data
- Generally not recommended
 - **DTO projection is better** / easier solution

Hibernate Documentation
calls it a “Marketing Feature”

Instead use DTO Projection

- Already discussed during **SELECT new Object**
 - Can select (project) only the properties you need
 - Better than lazy-loading properties!

```
TypedQuery<Home> query = em.createQuery(  
    "select new hibernate06.Home(p, a) "  
    + "from Person p " + "join p.address a ", Home.class);  
List<Home> homes = query.getResultList();
```

```
Person p = null;  
Address a = null;  
for (Home home : homes) {  
    p = home.getPerson();  
    a = home.getAddress();  
  
    System.out.println(p.getFirstName()  
        + " " + p.getLastName()  
        + " has a home in " + a.getCity());  
}
```

```
public class Home {  
    private Person person;  
    private Address address;  
  
    public Home(Person p, Address a) {  
        this.person = p;  
        this.address = a;  
    }  
}
```

Not an Entity
but a DTO class

Lazy and Eager Summary

- Lazy and Eager **specify WHEN not HOW**
 - When is generally not the problem
- It's good to know about these options
 - To understand how Hibernate works
 - While not the biggest source of problems or solutions



CS544 EA

Hibernate

Common Problems

Problems

- The essence of ORM related problems are that the database is being **overloaded**
 - Doesn't work properly anymore
 - Cannot handle the load
- Solution avenues are:
 - Lower the load by **using better techniques** (this chapter)
 - Spread the load by caching and scaling (we discuss caching)

Bad Queries

- The most common type of bad query is a **Cartesian Product**
 - **Caused by joining** 2 (or more) **collections**
 - Creates an ‘exploded’ resultset (takes the DB a long time)
- Hibernate will never generate such a query
 - Throws exception if 2 collections are set as eager on one Entity
 - But a (unaware) programmer can easily write such a query!

Code

Cust1 has 3 books and 3 movies
Cust2 1 book
Cust3 1 movie

```
Customer cust1 = new Customer("Frank", "Brown");
Customer cust2 = new Customer("Jane", "Terrien");
Customer cust3 = new Customer("John", "Doe");
cust1.addBook(new Book("Harry Potter and the Deathly Hallows"));
cust1.addBook(new Book("Unseen Academicals (Discworld)"));
cust1.addBook(new Book("The Color of Magic (Discworld)"));
cust1.addMovie(new Movie("Shrek"));
cust1.addMovie(new Movie("WALL-E"));
cust1.addMovie(new Movie("Howls Moving Castle"));
cust2.addBook(new Book("Twilight (The Twilight Saga, Book1)"));
cust3.addMovie(new Movie("Forgetting Sarah Marshall"));
em.persist(cust1);
em.persist(cust2);
em.persist(cust3);

em.getTransaction().commit();
em.clear();

em.getTransaction().begin();
TypedQuery<Customer> query = em.createQuery(
    "select c from Customer c left join c.movies left join c.books",
    Customer.class);
List<Customer> customers = query.getResultList();

em.getTransaction().commit();
```

Joining 2 collections

```
select
    customer0_.id as id1_3_0_,
    movies1_.id as id1_4_1_,
    books2_.id as id1_2_2_,
    customer0_.address_id as address_4_3_0_,
    customer0_.firstName as firstNam2_3_0_,
    customer0_.lastName as lastName3_3_0_,
    customer0_.salesRep_id as salesRep5_3_0_,
    movies1_.name as name2_4_1_,
    books2_.author_id as author_i3_2_2_,
    books2_.name as name2_2_2_
from
    Customer customer0_
left outer join
    Movie movies1_
        on customer0_.id=movies1_.movies_id
left outer join
    Book books2_
        on customer0_.id=books2_.books_id
```



Resultset

- Joining 2 collections **creates R x N x M rows**
 - R normal rows, N size of clct. 1, M size of clct. 2

FIRSTNAME0_0_	LASTNAME0_0_	TITLE1_1_	TITLE2_2_
Frank	Brown	Unseen Academicals (Discworld)	WALL-E
Frank	Brown	Unseen Academicals (Discworld)	Shrek
Frank	Brown	Unseen Academicals (Discworld)	Howls Moving Castle
Frank	Brown	The Color of Magic (Discworld)	WALL-E
Frank	Brown	The Color of Magic (Discworld)	Shrek
Frank	Brown	The Color of Magic (Discworld)	Howls Moving Castle
Frank	Brown	Harry Potter and the Deathly Hallows	WALL-E
Frank	Brown	Harry Potter and the Deathly Hallows	Shrek
Frank	Brown	Harry Potter and the Deathly Hallows	Howls Moving Castle
Jane	Terrien	Twilight (The Twilight Saga, Book1)	[null]
John	Doe	[null]	Forgetting Sarah Marshall

Redundancy

Very Inefficient!

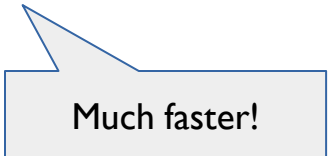
27 cells to give
7 pieces of data

Frank Brown ✓	Discworld ✓	Pixar ✓
Frank Brown	Discworld	Dream Works ✓
Frank Brown	Discworld	Studio Ghibli ✓
Frank Brown	Harry Potter ✓	Pixar
Frank Brown	Harry Potter	Dream Works
Frank Brown	Harry Potter	Studio Ghibli
Frank Brown	Twilight ✓	Pixar
Frank Brown	Twilight	Dream Works
Frank Brown	Twilight	Studio Ghibli

Model © Prof. Rene de Jong

N + 1 Problem

- The N+1 problem is where Hibernate executes **many small selects** to load related data
 - This data could have been loaded in **one big select**
- People sometimes associate it with lazy loading
 - But happens with eager loading too!
 - It's just Hibernate not knowing **how** to best load data



Much faster!

Lazy Collections N+1

- By default Hibernate lazily loads collections
 - A good default, they can contain a lot of data
- If we create **a query** for all SalesReps
 - Then **use a loop** to get the customers of those reps
 - I select for the salesreps (say there are 10)
 - 10 selects, one for each collection of customers

Code


```
em.getTransaction().begin();
```

```
SalesRep sr1 = new SalesRep("John Willis");  
SalesRep sr2 = new SalesRep("Mary Long");
```

```
sr1.addCustomer(new Customer("Frank", "Brown"));  
sr1.addCustomer(new Customer("Jane", "Terrien"));  
sr2.addCustomer(new Customer("John", "Doe"));  
sr2.addCustomer(new Customer("Carol", "Reno"));
```

```
em.persist(sr1);  
em.persist(sr2);  
em.getTransaction().commit();
```

```
TypedQuery<SalesRep> query =  
    em.createQuery("from SalesRep", SalesRep.class);  
List<SalesRep> salesReps = query.getResultList();  
for(SalesRep s : salesReps) {  
    System.out.println(s.getCustomers().get(0));  
}
```



I select for the salesreps
N selects for each
list of customers

Hibernate:

```
select  
    salesrep0_.id as id1_1_,  
    salesrep0_.name as name2_1_  
from  
    SalesRep salesrep0_
```

Hibernate:

```
select  
    customers0_.salesRep_id as salesRep4_0_0_,  
    customers0_.id as id1_0_0_,  
    customers0_.id as id1_0_1_,  
    customers0_.firstName as firstNam2_0_1_,  
    customers0_.lastName as lastName3_0_1_,  
    customers0_.salesRep_id as salesRep4_0_1_  
from  
    Customer customers0_  
where  
    customers0_.salesRep_id=?
```

Hibernate:

```
select  
    customers0_.salesRep_id as salesRep4_0_0_,  
    customers0_.id as id1_0_0_,  
    customers0_.id as id1_0_1_,  
    customers0_.firstName as firstNam2_0_1_,  
    customers0_.lastName as lastName3_0_1_,  
    customers0_.salesRep_id as salesRep4_0_1_  
from  
    Customer customers0_  
where  
    customers0_.salesRep_id=?
```

Visually

Select * from Salesrep

I select

N selects, when accessing the collections

Salesrep #1

Select * from Customer
where salesrep_id = 1

Customer #1
Customer #2
Customer #3

Salesrep #2

Select * from Customer
where salesrep_id = 2

Customer #4
Customer #5
Customer #6

Salesrep #3

Select * from Customer
where salesrep_id = 3

Customer #7
Customer #8
Customer #9

Salesrep #4

Select * from Customer
where salesrep_id = 4

Customer #10
Customer #11
Customer #12

Salesrep #5

Select * from Customer
where salesrep_id = 5

Customer #13
Customer #14
Customer #15

Size of
collection
does not
matter

Eager References N+1

- By default Hibernate uses eager loading for

- @OneToOne and @ManyToOne
- If eager associations are not yet fulfilled
 - Hibernate will execute select statements to fix it

Good policy.
Cost of joining a single row is low, and generally reduces selects

- If you execute **1 query** for all customers

References are Eager by default

- Without Join Fetch-ing the @ManyToOne SalesRep
- Hibernate will 'fix' this right away with **N extra selects**

Doesn't even need a loop

Code

Each customer has its own Rep

I select for customers,
N selects for the reps

```
em.getTransaction().begin();
Customer cust1 = new Customer("Frank", "Brown");
Customer cust2 = new Customer("Jane", "Terrien");
Customer cust3 = new Customer("John", "Doe");
Customer cust4 = new Customer("Carol", "Reno");
cust1.setSalesRep(new SalesRep("John Willis"));
cust2.setSalesRep(new SalesRep("Mary Long"));
cust3.setSalesRep(new SalesRep("Ted Walker"));
cust4.setSalesRep(new SalesRep("Keith Rogers"));
```

```
em.persist(cust1);
em.persist(cust2);
em.persist(cust3);
em.persist(cust4);
em.getTransaction().commit();
```

```
List<Customer> customers = em.createQuery(
    "from Customer").getResultList();
```

No loop or anything.

Hibernate executes the selects
to fix the missing eager references

```
Hibernate:
select
  customer0_.id as id1_0_,
  customer0_.firstName as firstNam2_0_,
  customer0_.lastName as lastName3_0_,
  customer0_.salesRep_id as salesRep4_0_
from
  Customer customer0_
```

```
Hibernate:
select
  salesrep0_.id as id1_1_0_,
  salesrep0_.name as name2_1_0_
from
  SalesRep salesrep0_
where
  salesrep0_.id=?
```

```
Hibernate:
select
  salesrep0_.id as id1_1_0_,
  salesrep0_.name as name2_1_0_
from
  SalesRep salesrep0_
where
  salesrep0_.id=?
```

```
Hibernate:
select
  salesrep0_.id as id1_1_0_,
  salesrep0_.name as name2_1_0_
from
  SalesRep salesrep0_
where
  salesrep0_.id=?
```

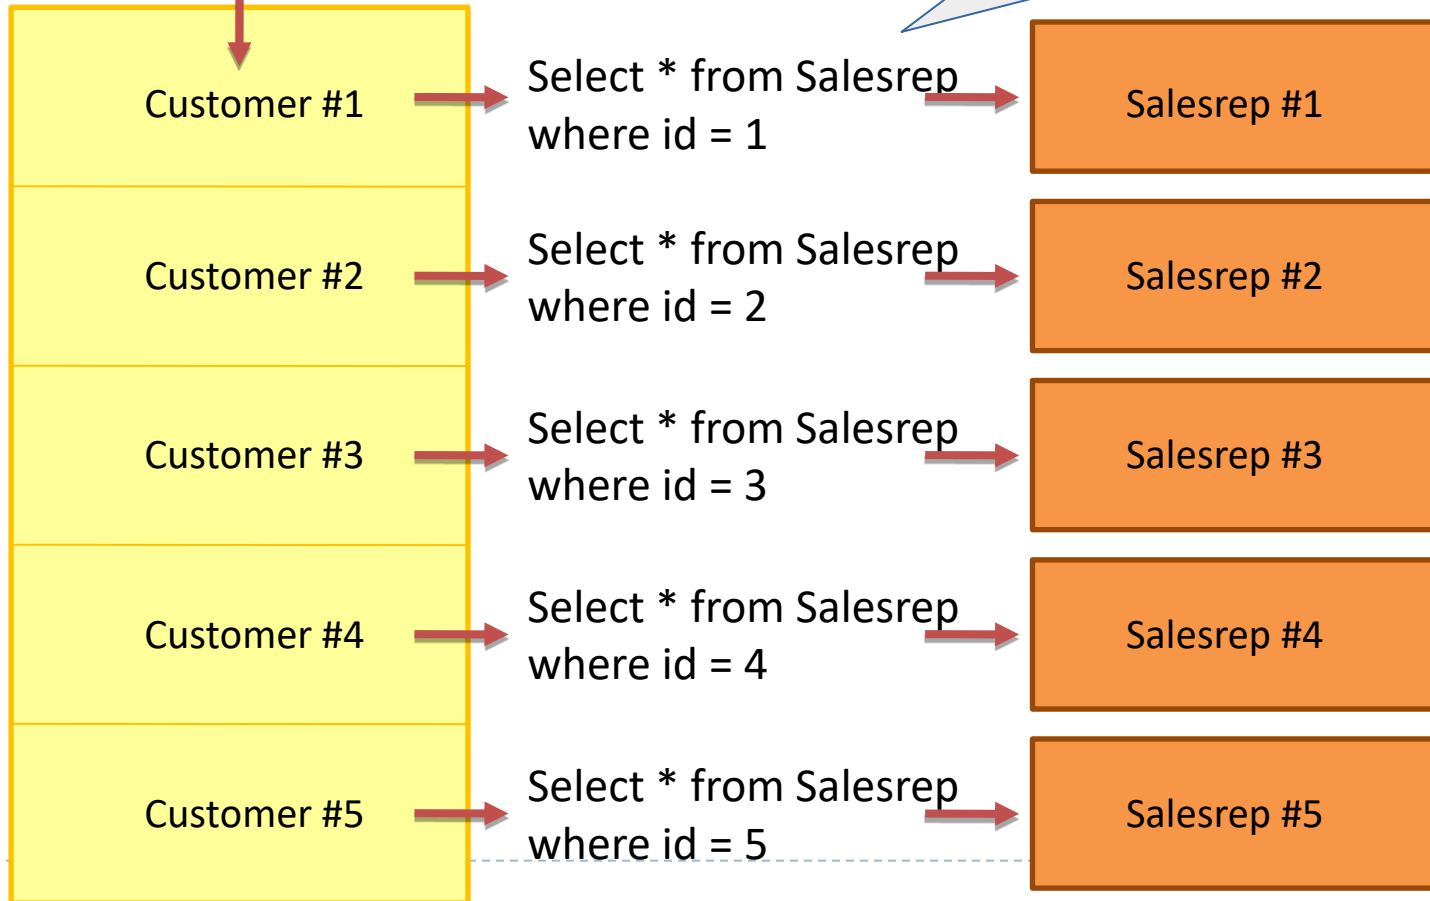
Hibernate:

Visually

Select * from Customer

I select

N selects, right away




Changing Doesn't Help

- Changing the references to LAZY

- Just makes it so that Hibernate doesn't load the entities until you access them (with a loop)

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private Long id;
    private String firstName;
    private String lastName;
    @ManyToOne(fetch=FetchType.LAZY
)
    private SalesRep salesRep;
```



- Similarly, changing the collection to EAGER

- Makes the N selects happen right away
- The **problem is not in WHEN, but HOW**

Solutions

- The solution for the **Cartesian product** is simple:
 - **Don't join 2 or more** collections in one query
 - Join max 1, use separate queries for the others
 - Similarly other bad queries can be analyzed and fixed
- The solution for **N+1 is not that easy**
 - We'll look at potential strategies coming up



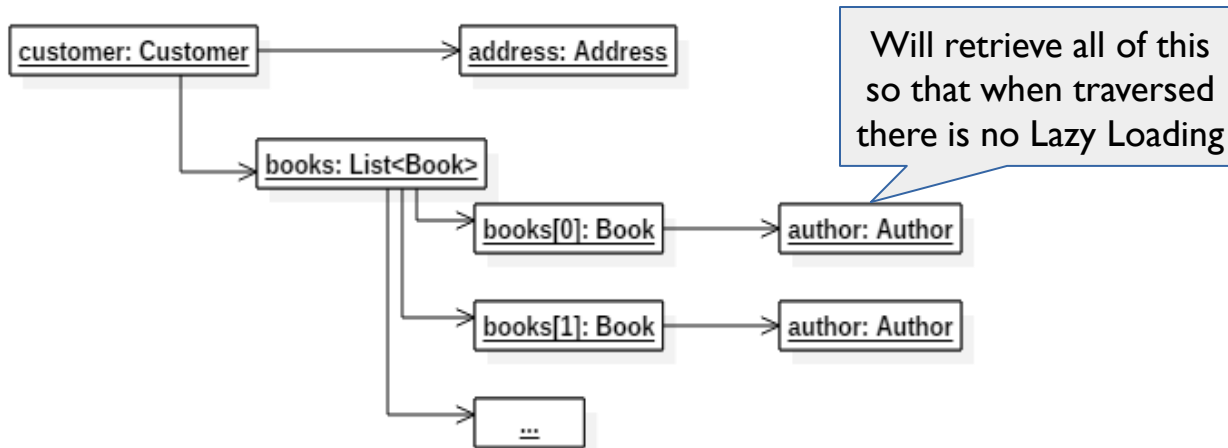
CS544 EA

Hibernate

Optimization: Entity Graph

Entity Graph

- Added in JPA 2.1 (most recent)
 - Specify a **Graph of connected Entities** to retrieve
 - Example: When retrieving a customer we also want to get his address, and all the books he bought, and the author of each of those books



Domain

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private Long id;
    private String firstName;
    private String lastName;
    @OneToOne(cascade=CascadeType.ALL)
    private Address address;
    @OneToMany(cascade=CascadeType.ALL)
    @JoinColumn
    private List<Book> books = new ArrayList<>();
}
```

```
@Entity
public class Book {
    @Id
    @GeneratedValue
    private Long id;
    private String name;
    @OneToOne(cascade=CascadeType.ALL)
    private Author author;
}
```

```
Customer cust1 = new Customer("Frank", "Brown");
Customer cust2 = new Customer("Jane", "Terrien");
Customer cust3 = new Customer("John", "Doe");

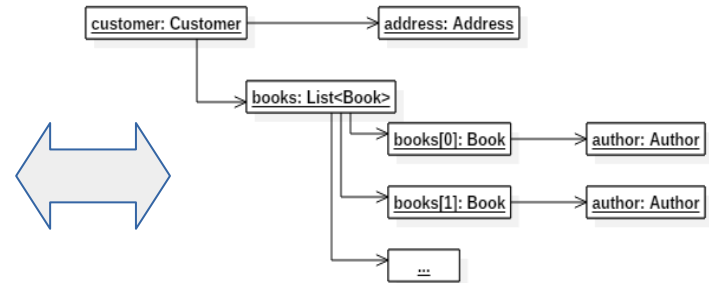
cust1.addBook(
    new Book("Harry Potter and the Deathly Hallows",
        new Author("J.K. Rowlings")));
cust1.addBook(
    new Book("Unseen Academicals (Discworld)",
        new Author("Terry Pratchett")));
cust1.addBook(
    new Book("The Color of Magic (Discworld)",
        new Author("Terry Pratchett")));
cust2.addBook(
    new Book("Twilight (The Twilight Saga, Book1)",
        new Author("Stephenie Meyer")));
cust1.setAddress(new Address("Fairfield", "Iowa"));
cust2.setAddress(new Address("Chicago", "Illinois"));
em.persist(cust1);
em.persist(cust2);
em.persist(cust3);
```

```
@Entity
public class Address {
    @Id
    @GeneratedValue
    private Long id;
    private String city;
    private String state;
}
```

EntityGraph

- The purpose of the entity graph is:
 - To **indicate which** references should change to **load eagerly** (in a query or .find())
 - AttributeNodes specify attributes / references
 - SubGraph can be used to go into other Entities

```
EntityGraph<Customer> graph =  
    em.createEntityGraph(Customer.class);  
graph.addAttributeNodes("address");  
graph.addSubgraph("books").addAttributeNodes("author");
```



.createQuery()

```
EntityGraph<Customer> graph =  
em.createEntityGraph(Customer.class);  
graph.addAttributeNodes("address");  
graph.addSubgraph("books").addAttributeNodes("author");
```

```
TypedQuery<Customer> query = em.createQuery(  
    "from Customer where firstName like :name",  
    Customer.class);  
query.setParameter("name", "J%");  
query.setHint("javax.persistence.fetchgraph", graph);
```

```
List<Customer> customers = query.getResultList();  
System.out.println(customers.size());
```

The EntityGraph is passed
as a query Hint

Hibernate loads the entire graph into cache
While returning the Customer as query result

Hibernate:

```
select  
    customer0_.id as id1_3_0_,  
    customer0_.address_id as address_4_3_0_,  
    customer0_.firstName as firstNam2_3_0_,  
    customer0_.lastName as lastName3_3_0_,  
    address1_.id as id1_0_1_,  
    address1_.city as city2_0_1_,  
    address1_.state as state3_0_1_,  
    books2_.books_id as books_id4_2_2_,  
    books2_.id as id1_2_2_,  
    books2_.id as id1_2_3_,  
    books2_.author_id as author_i3_2_3_,  
    books2_.name as name2_2_3_,  
    author3_.id as id1_1_4_,  
    author3_.name as name2_1_4_  
from  
    Customer customer0_  
left outer join  
    Address address1_  
        on customer0_.address_id=address1_.id  
left outer join  
    Book books2_  
        on customer0_.id=books2_.books_id  
left outer join  
    Author author3_  
        on books2_.author_id=author3_.id  
where  
    customer0_.id=?
```

2

.find()

```
EntityGraph<Customer> graph
    = em.createEntityGraph(Customer.class);
graph.addAttributeNodes("address");
graph.addSubgraph("books").addAttributeNodes("author");

Map<String, Object> properties = new HashMap<>();
properties.put("javax.persistence.fetchgraph", graph);

em.find(Customer.class, 1L, properties);
```

Hints are passed as
properties Map to .find()

Hibernate loads the entire graph into cache
giving us the Root (Customer) entity

Hibernate:

```
select
    customer0_.id as id1_3_0_,
    customer0_.address_id as address_4_3_0_,
    customer0_.firstName as firstNam2_3_0_,
    customer0_.lastName as lastName3_3_0_,
    address1_.id as id1_0_1_,
    address1_.city as city2_0_1_,
    address1_.state as state3_0_1_,
    books2_.books_id as books_id4_2_2_,
    books2_.id as id1_2_2_,
    books2_.id as id1_2_3_,
    books2_.author_id as author_i3_2_3_,
    books2_.name as name2_2_3_,
    author3_.id as id1_1_4_,
    author3_.name as name2_1_4_
from
    Customer customer0_
left outer join
    Address address1_
        on customer0_.address_id=address1_.id
left outer join
    Book books2_
        on customer0_.id=books2_.books_id
left outer join
    Author author3_
        on books2_.author_id=author3_.id
where
    customer0_.id=?
```

Entity Graph and N+1

- An entity graph can be a solution to N+1
 - Load all the needed entities in one query
- Potential problems:
 - You can not make more than one collection eager
 - Eager associations from your graph / result to other entities(outside your graph) **still cause N+1** (see eager references N+1)

If you have references going out
make sure they're lazy!



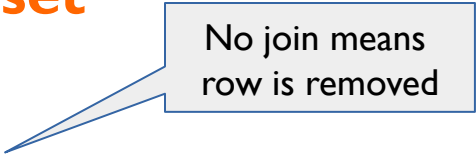
CS544 EA

Hibernate

Optimization: Join Fetch

Join Fetch

- Before Fetch Graphs were added to JPA
 - Queries could already do “Join Fetch”
- Like EntityGraph Join Fetch-ed entities are:
 - **Added to the cache**
 - **Not added to the result set**
- Unlike EntityGraph
 - Join Fetch now uses inner join (instead of left outer)



No join means
row is removed

Join Fetch

No SELECT clause needed
joined entities not added to ResultSet

```
TypedQuery<Customer> query = em.createQuery(
    "from Customer c "
    + "JOIN FETCH c.address a "
    + "JOIN FETCH c.books b "
    + "JOIN FETCH b.author "
    + "WHERE c.firstName like :name", Customer.class);
query.setParameter("name", "J%");
List<Customer> customers = query.getResultList();
System.out.println(customers.size());
```

Remember: don't join (fetch)
multiple collections!

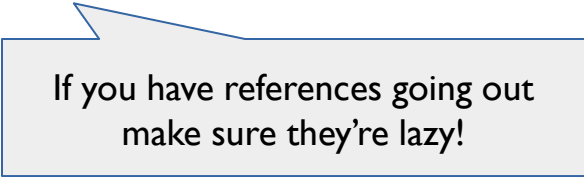
Hibernate:

```
select
  customer0_.id as id1_3_0_,
  address1_.id as id1_0_1_,
  books2_.id as id1_2_2_,
  author3_.id as id1_1_3_,
  customer0_.address_id as address_4_3_0_,
  customer0_.firstName as firstNam2_3_0_,
  customer0_.lastName as lastName3_3_0_,
  address1_.city as city2_0_1_,
  address1_.state as state3_0_1_,
  books2_.author_id as author_i3_2_2_,
  books2_.name as name2_2_2_,
  books2_.books_id as books_id4_2_0_,
  books2_.id as id1_2_0_,
  author3_.name as name2_1_3_
from
  Customer customer0_
inner join
  Address address1_
    on customer0_.address_id=address1_.id
inner join
  Book books2_
    on customer0_.id=books2_.books_id
inner join
  Author author3_
    on books2_.author_id=author3_.id
where
  customer0_.firstName like ?
```

1

Join Fetch and N+1

- Join Fetch can be a solution for N+1
 - Load all the needed objects in one query
- Same potential problems:
 - You can not Join Fetch more than one collection eager
 - Eager associations from your graph / result to other entities **still cause N+1** (see eager references N+1)



If you have references going out
make sure they're lazy!



CS544 EA

Hibernate

Optimization: BatchSize

@BatchSize

- Hibernate extension commonly used for N+1
 - Helps by loading **several collections in one 'batch'**
 - Gets another batch when previous batch is empty
 - Turns $N + 1$ into: $\text{ceil}(N / \text{Batchsize}) + 1$

```
import org.hibernate.annotations.BatchSize;
```

```
@Entity
```

```
public class SalesRep {
```

```
    @Id
```

```
    @GeneratedValue
```

```
    private Long id;
```

```
    private String name;
```

```
    @OneToMany(mappedBy = "salesRep", cascade=CascadeType.ALL)
```

```
    @BatchSize(size=3)
```

```
    private List<Customer> customers = new ArrayList<>();
```

Typical size between
3 and 15

I select

Select * from Salesrep

N / Batchsize
selects

Each select returns
a 'Batch' of collections

Salesrep #1

Select * from Customer
where salesrep_id in (1, 2, 3)

Customer #1
Customer #2
Customer #3

Salesrep #2

Works because it knows
the IDs of next SalesReps

Customer #4
Customer #5
Customer #6

Salesrep #3

Customer #7
Customer #8
Customer #9

Salesrep #4

Select * from Customer
where salesrep_id in (4, 5)

Customer #10
Customer #11
Customer #12

Salesrep #5

Customer #13
Customer #14
Customer #15

Code

```
em.getTransaction().begin();
```

```
SalesRep sr1 = new SalesRep("John Willis");  
SalesRep sr2 = new SalesRep("Mary Long");
```

```
sr1.addCustomer(new Customer("Frank", "Brown"));  
sr1.addCustomer(new Customer("Jane", "Terrien"));  
sr2.addCustomer(new Customer("John", "Doe"));  
sr2.addCustomer(new Customer("Carol", "Reno"));
```

```
em.persist(sr1);  
em.persist(sr2);  
em.getTransaction().commit();
```

```
System.out.println("Before query");  
TypedQuery<SalesRep> query =  
    em.createQuery("from SalesRep", SalesRep.class);  
List<SalesRep> salesReps = query.getResultList();  
System.out.println("After query");  
for (SalesRep s : salesReps) {  
    s.getCustomers().get(0).getFirstName();  
}  
System.out.println("After loop");
```

Before query

Hibernate:

```
select  
    salesrep0_.id as id1_3_,  
    salesrep0_.name as name2_3_  
from  
    SalesRep salesrep0_
```

After query

Hibernate:

```
select  
    customers0_.salesRep_id as salesRep4_1_1_,  
    customers0_.id as id1_1_1_,  
    customers0_.id as id1_1_0_,  
    customers0_.firstName as firstNam2_1_0_,  
    customers0_.lastName as lastName3_1_0_,  
    customers0_.salesRep_id as salesRep4_1_0_  
from  
    Customer customers0_  
where  
    customers0_.salesRep_id in (  
        ?, ?  
    )
```

After loop

Batch loaded when
customer first needed

BatchSize and N+1

- @BatchSize does not completely eliminate N+1
 - It does significantly reduce it
 - N+1 becomes: $\lceil N / \text{batchsize} \rceil + 1$
- Potential Problems:
 - N+1 effects not completely removed (just reduced)
 - Static, always on, no way to not use it
 - May load (batchsize - 1) objects not needed



CS544 EA

Hibernate

Optimization: FetchMode.SUBSELECT

FetchMode.SUBSELECT

- Hibernate Extension commonly used for N+1
 - Turns **N + 1** into **1 + 1**
 - Can load **too much data**

```
import org.hibernate.annotations.Fetch;
import org.hibernate.annotations.FetchMode;

@Entity
public class SalesRep {
    @Id
    @GeneratedValue
    private Long id;
    private String name;
    @OneToMany(mappedBy = "salesRep", cascade=CascadeType.ALL)
    @Fetch(FetchMode.SUBSELECT)
    private List<Customer> customers = new ArrayList<>();
}
```

@Fetch

I select

Select * from Salesrep

Salesrep #1

Salesrep #2

Salesrep #3

Salesrep #4

Salesrep #5

I select

**Select * from Customer
where salesrep_id in
(Select id from Salesrep)**

Uses a sub-select
to specify IDs

Select returns
all Customers related to
our SalesReps

Customer #1
Customer #2
Customer #3

Customer #4
Customer #5
Customer #6

Customer #7
Customer #8
Customer #9

Customer #10
Customer #11
Customer #12

Customer #13
Customer #14
Customer #15

Code

```
em.getTransaction().begin();
```

```
SalesRep sr1 = new SalesRep("John Willis");  
SalesRep sr2 = new SalesRep("Mary Long");
```

```
sr1.addCustomer(new Customer("Frank", "Brown"));  
sr1.addCustomer(new Customer("Jane", "Terrien"));  
sr2.addCustomer(new Customer("John", "Doe"));  
sr2.addCustomer(new Customer("Carol", "Reno"));
```

```
em.persist(sr1);  
em.persist(sr2);  
em.getTransaction().commit();
```

```
System.out.println("Before query");  
TypedQuery<SalesRep> query =  
    em.createQuery("from SalesRep where id < 1000",  
        SalesRep.class);  
List<SalesRep> salesReps = query.getResultList();  
System.out.println("After query");  
for (SalesRep s : salesReps) {  
    s.getCustomers().get(0).getFirstName();  
}  
System.out.println("After loop");
```

Before query

Hibernate:

```
select  
    salesrep0_.id as id1_3_,  
    salesrep0_.name as name2_3_  
from  
    SalesRep salesrep0_  
where  
    salesrep0_.id<1000
```

After query

Hibernate:

```
select  
    customers0_.salesRep_id as salesRep4_1_1_,  
    customers0_.id as id1_1_1_,  
    customers0_.id as id1_1_0_,  
    customers0_.firstName as firstNam2_1_0_,  
    customers0_.lastName as lastName3_1_0_,  
    customers0_.salesRep_id as salesRep4_1_0_  
from  
    Customer customers0_  
where  
    customers0_.salesRep_id in (  
        select  
            salesrep0_.id  
        from  
            SalesRep salesrep0_  
        where  
            salesrep0_.id<1000  
    )
```

After loop

SubSelect Query
executes when
first customer needed

Includes
where

SubSelect and N+1

- SUBSELECT is like a **supercharged BatchSize**
 - Works on same principle but goes all the way
 - Turns $N + 1$ into $1 + 1$
- Potential problem:
 - Static, always on, no way to not use it
 - May load everything even if we only needed 1

N+1 Summary


- Entity Graph & Join Fetch (very similar)
 - Dynamic, you can chose when to use them
 - Never join more than 1 collection
 - Be careful with eager associations outside result
- BatchSize
 - Static → may load a bit too much
- SubSelect
 - Static → may load way too much

While I personally like
BatchSize and Join Fetch
you really need to analyze
what works best
for your situation

Cartesian Product Summary

- ▶ To fix a Cartesian product, use more queries.

- ▶ For example if we need to make a report on which active Salesrep sold the most books or movies



- ▶ Create the following queries:
 - ▶ From Salesrep s join s.customers c where s.active = 1
 - ▶ From Customer c join c.books b where c.salesrep.active=1
 - ▶ From Customer c join c.movies m where c.salesrep.active=1
- ▶ Once data is loaded you can use it (without N+1 fear)





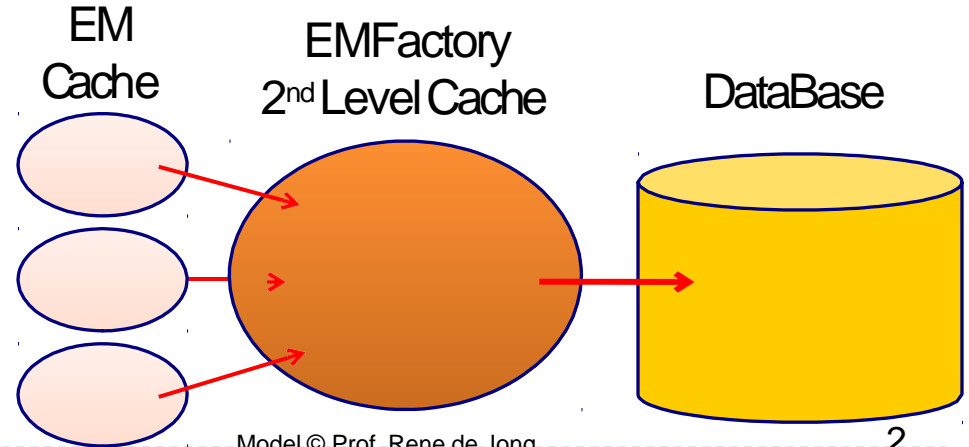

CS544 EA

Hibernate

Optimization: 2nd Level Cache


2nd Level Caching

- By default JPA only uses EntityManager cache
 - Very short term cache
- To reduce hits on the DB
 - Objects can also be cached for longer
 - Managed by EntityManagerFactory
 - Shared by all EntityManagers



Caching VS Optimization

- Caching can be seen as a **form of scaling**
 - Doesn't solve bad queries
 - But can alleviate pressure on the DB
- Caching is a large and interesting field
 - We will look at some basics
 - Be aware that **improper configuration** can create situations that are **hard to debug** (cached versions != DB versions)



What to cache?

- Good candidates for caching:
 - **Do not change**, or change rarely
 - Are modified only by your app
 - Are non-critical to the app
- Typically: **Reference data**

4 Caching Strategies

Stricter and therefore Slower

- **Read Only**: very fast strategy, but can only be used for data that never changes
- **Non-Strict Read-Write**: data may be stale for a while, but gets refreshed at a timeout
- **Read-Write**: prevents stale data, but at a cost. Use for read-mostly data in a non-clustered setup
- **Transactional**: Can prevent stale data in a clustered environment. Can be used for read-mostly data

Cache Providers

- Hibernate can have **only one** provider per EMF

Provider	Read Only	Non Strict Read Write	Read Write	Transactional
EHCache	✓	✓	✓	
OSCache	✓	✓	✓	
SwarmCache	✓	✓		
JBoss Cache 1.x	✓			✓
JBoss Cache 2.x	✓			✓

Annotate Classes with Strategy

- Using Hibernate's **@Cache** annotation

```
@Entity
@Cache(usage= CacheConcurrencyStrategy.READ_ONLY)
public class Customer implements Serializable {
    @Id
    @GeneratedValue
    private Long id;
    private String firstName;
    private String lastName;

    public Customer() {
    }

    public Customer(String firstName, String lastName) {
        this.firstName = firstName;
        this.lastName = lastName;
    }
}
```

Setup Cache Provider

- Inside **persistence.xml**

```
<properties>
  <!-- Useful for analyzing problems -->
  <property name="hibernate.show_sql" value="true"/>
  <property name="hibernate.format_sql" value="false"/>

  <!-- 2nd Level Caching -->
  <!--By default is true, second-level-cache is enabled-->
  <property name="hibernate.cache.use_second_level_cache" value="true"/>
  <!--enable query caching-->
  <!--<property name="hibernate.cache.use_query_cache" value="true" />-->
  <property name="hibernate.cache.region.factory_class"
    value="org.hibernate.cache.jcache.JCacheRegionFactory"/>
  <property name="net.sf.ehcache.configurationResourceName" value="/ehcache.xml"/>
  <property name="hibernate.javax.cache.provider" value="org.ehcache.jsr107.EhcacheCachingProvider"/>
  <!-- To analyze cache performance -->
  <property name="hibernate.generate_statistics" value="true"/>

  ...
</properties>
```

Configure Cache Provider

```
<ehcache>
  <diskStore path="java.io.tmpdir"/>
  <defaultCache
    maxElementsInMemory="10000"
    eternal="false"
    timeToIdleSeconds="120"
    timeToLiveSeconds="120"
    overflowToDisk="true" />
```

General Config

```
<cache name="cacheDemo.Category"
  maxElementsInMemory="50"
  eternal="true"
  timeToIdleSeconds="0"
  timeToLiveSeconds="0"
  overflowToDisk="false" />
```

Config for an Entity

```
<cache
name="cacheDemo.Category.customers"
  maxElementsInMemory="50"
  eternal="false"
  timeToIdleSeconds="3600"
  timeToLiveSeconds="7200"
  overflowToDisk="false" />
```

Config for a Collection

```
<cache name="cacheDemo.SalesRep"
  maxElementsInMemory="500"
  eternal="false"
  timeToIdleSeconds="1800"
  timeToLiveSeconds="10800"
  overflowToDisk="false" />
```

```
</ehcache>
```

Statistics

```
SessionFactory sessionFactory = emf.unwrap(SessionFactory.class);
```

```
Statistics stats = sessionFactory.getStatistics();
```

```
long hits = stats.getSecondLevelCacheHitCount();
```

```
long misses = stats.getSecondLevelCacheMissCount();
```

```
long puts = stats.getSecondLevelCachePutCount();
```

```
System.out.printf("\nGeneral 2nd Level Cache Stats\n");
```

```
System.out.printf("Hit: %d Miss: %d Put: %d\n", hits, misses, puts);
```

General 2nd level
cache statistics

```
org.hibernate.stat.CacheRegionStatistics customerCacheStats =
```

```
    stats.getCacheRegionStatistics("hibernate07.Customer");
```

```
long srCurrent = customerCacheStats.getElementCountInMemory();
```

```
long srMemsize = customerCacheStats.getSizeInMemory();
```

```
long srHits = customerCacheStats.getHitCount();
```

```
long srMisses = customerCacheStats.getMissCount();
```

```
long srPuts = customerCacheStats.getPutCount();
```

```
System.out.printf("\ncustomerCache Cache Region - Size: %d Holds: %d\n", srMemsize, srCurrent);
```

```
System.out.printf("Hit: %d Miss: %d Put: %d\n", srHits, srMisses, srPuts);
```

Statistics for a
specific cache region

```
SessionFactory sessionFactory = emf.unwrap(SessionFactory.class);
```

```
Statistics stats = sessionFactory.getStatistics();
```

```
Stats.clear();
```

```
stats.setStatisticsEnabled(true);
```

```
...
```

```
stats.setStatisticsEnabled(false);
```

You can also programmatically
turn stats on and off for
more targeted measuring

Summary

- Premature optimization is the root of all evil
- If your DB is under load use better techniques:
 - Check the SQL that's generated
 - Make sure no Cartesian Product
 - If you find N+1 problems use:
 - Dynamic: EntityGraph, or JoinFetch
 - Static: @BatchSize, or SubSelect
- Use Caching (or scaling) to further reduce load